

3. Process Concept

[ECE321/ITP302] Operating Systems

Objectives



- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore inter-process communication using shared memory and message passing
- To describe communication in client-server systems

Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

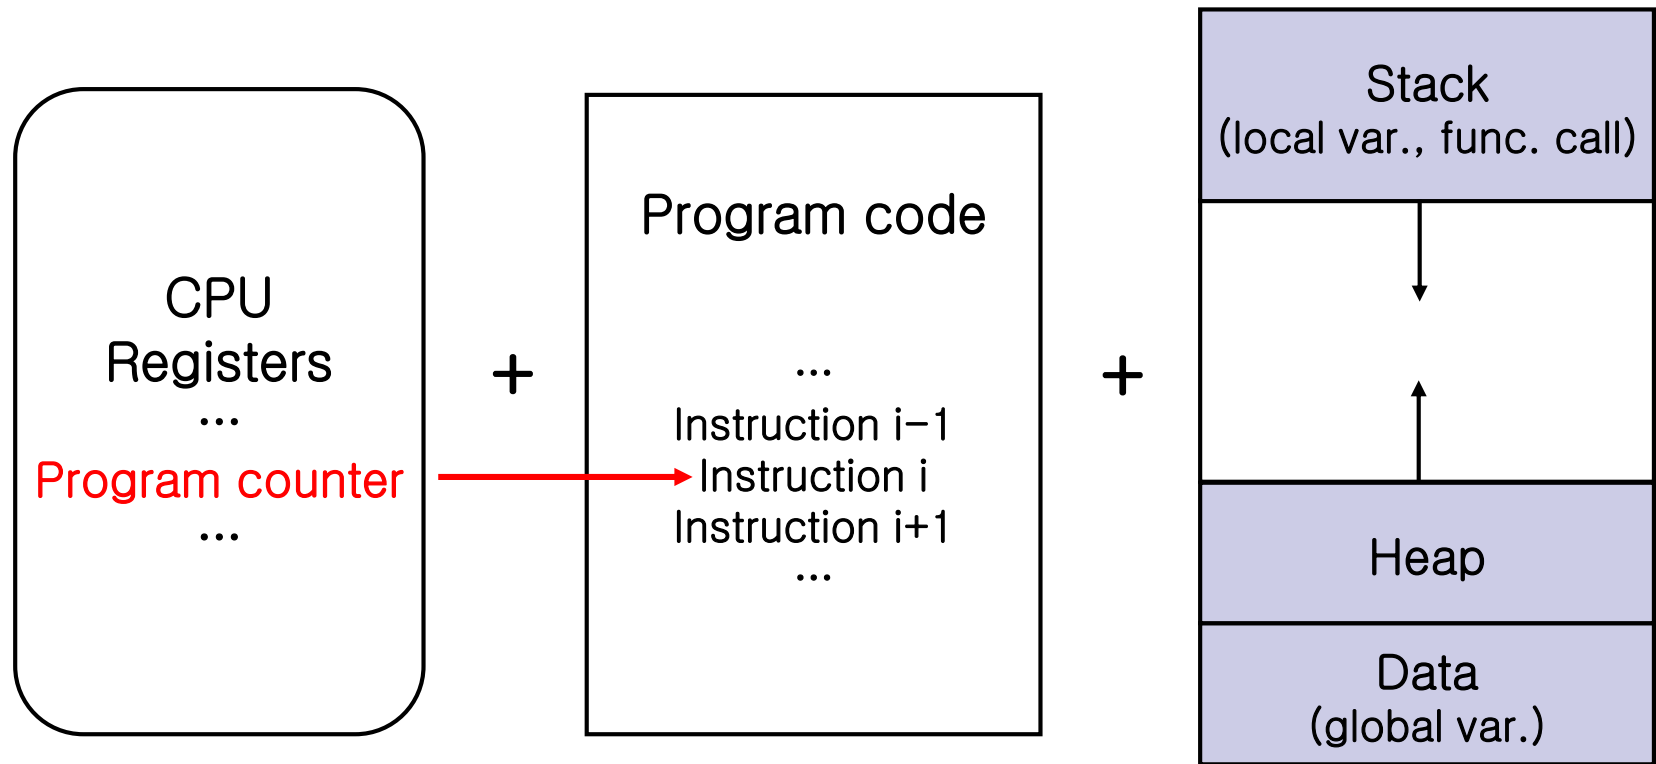
Process Concept



- An operating system executes a variety of programs:
 - Batch system – [jobs](#)
 - Time-shared systems – [user programs](#) or [tasks](#)
- Textbook uses the terms *job* and *process* almost interchangeably
- [Process](#) – a program in execution; process execution must progress in sequential fashion
- Multiple parts
 - The program code, also called [text section](#)
 - Current activity including [program counter](#), processor registers
 - [Stack](#) containing temporary data
 - Function parameters, return addresses, local variables
 - [Data section](#) containing global variables
 - [Heap](#) containing memory dynamically allocated during run time
- Program is *passive* entity stored on disk ([executable file](#)), process is *active*
 - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider multiple users executing the same program


Process

- **Process** = program in execution + resource

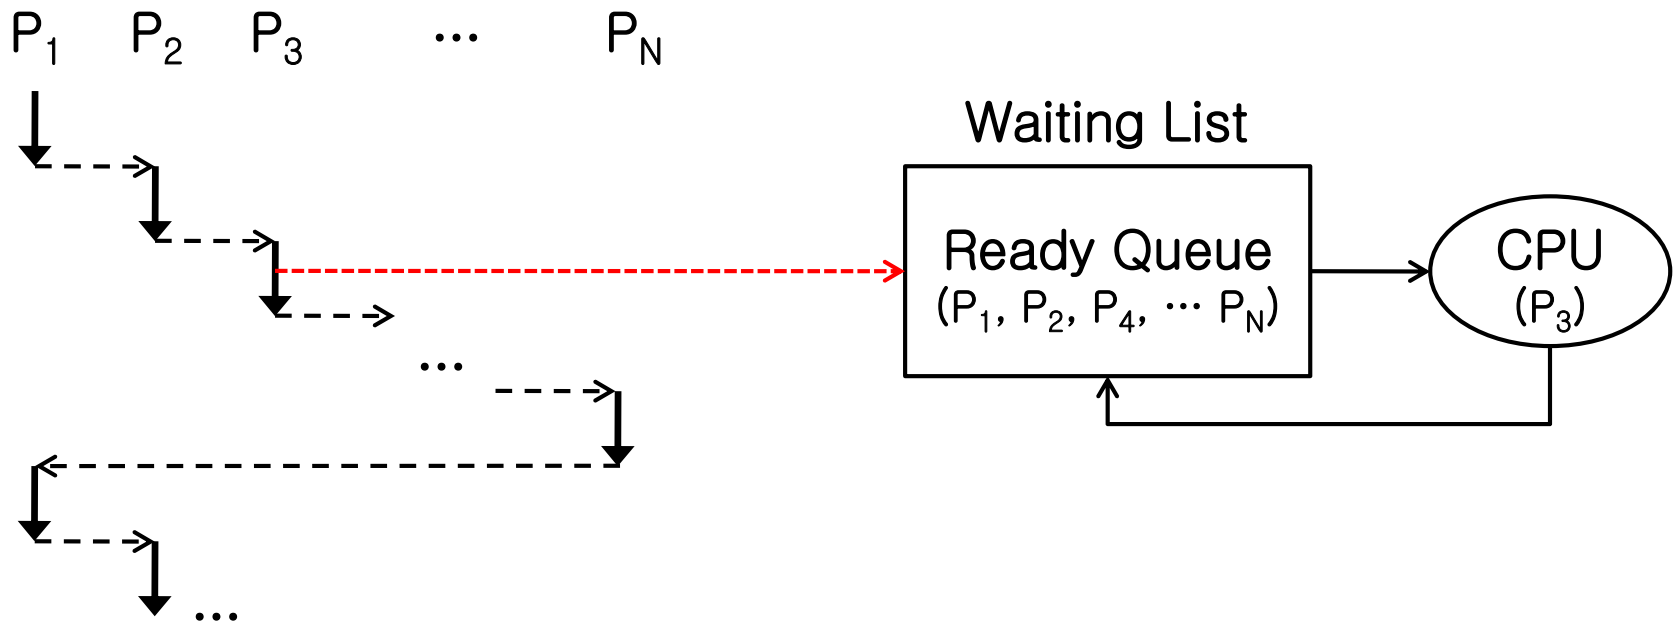


Process State

- **New:** being created
- **Running:** in execution
 - Only one process can be running on a processor at any time
- **Ready:** waiting to be assigned to a processor
- **Waiting:** waiting for some event to occur
- **Terminated**




Ready/Running State



Process Control Block (PCB)

- OS manages processes using PCB
 - Process Control Block (PCB)
repository for any information about process



Contents	Examples
Process state	new, ready, running, waiting, terminated, ...
Process number	pid (Process ID)
CPU Registers	<u>program counter</u> (address of next instruction to execute) accumulator, general registers, stack pointer, ...
CPU Scheduling info.	priority, pointer to queue, ...
Memory-management info.	base and limit registers, page/segment table, ...
Accounting info.	CPU-time used, time limits, account #, ...
I/O status info.	List of open files, I/O devices allocated

Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

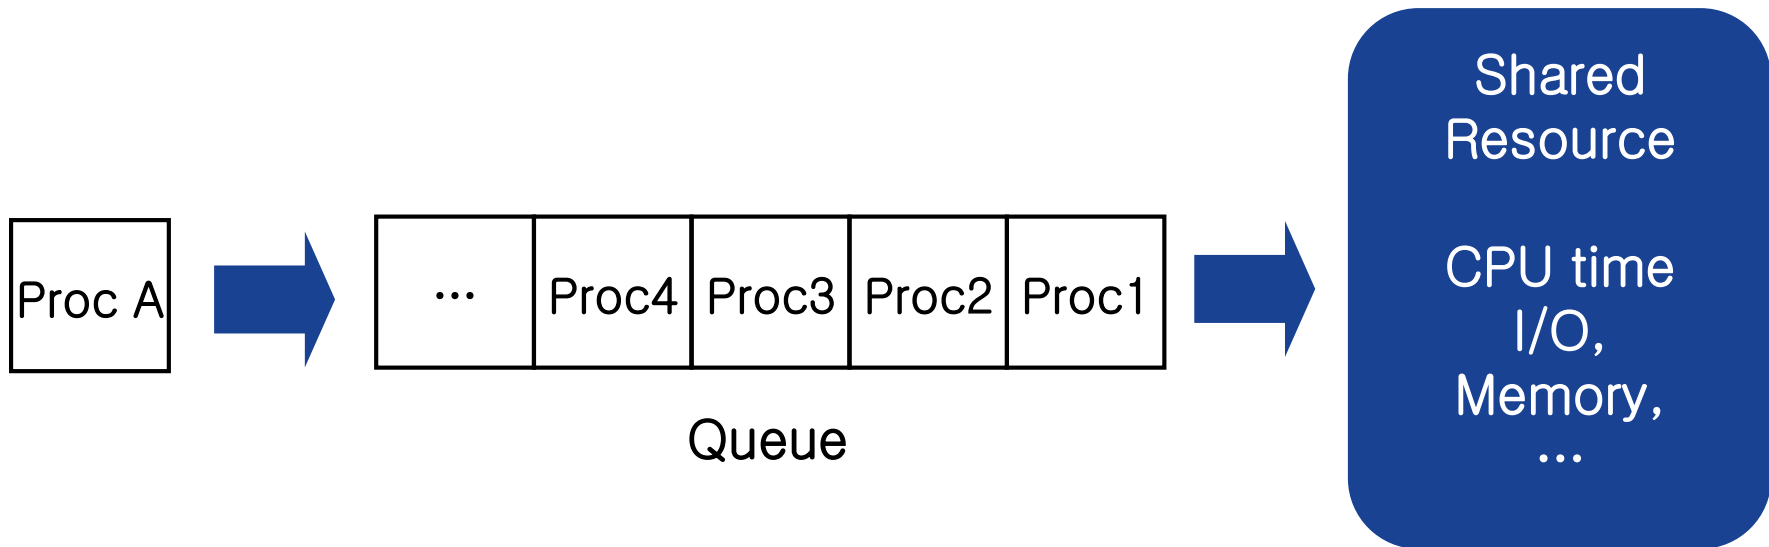
Process Scheduling



- **Scheduling**: assigning tasks to a set of resources
- **Process scheduling**: selecting a process to execute on CPU
 - Only one process can run on each processor at a time.
 - Other processes should wait
- **Objectives of scheduling**
 - Maximize CPU utilization (multiprogramming)
 - Users can interact with each program (time sharing)

Scheduling Queue

- **Scheduling queue**: waiting list of processes for CPU time or other resources



Types of Scheduling Queues



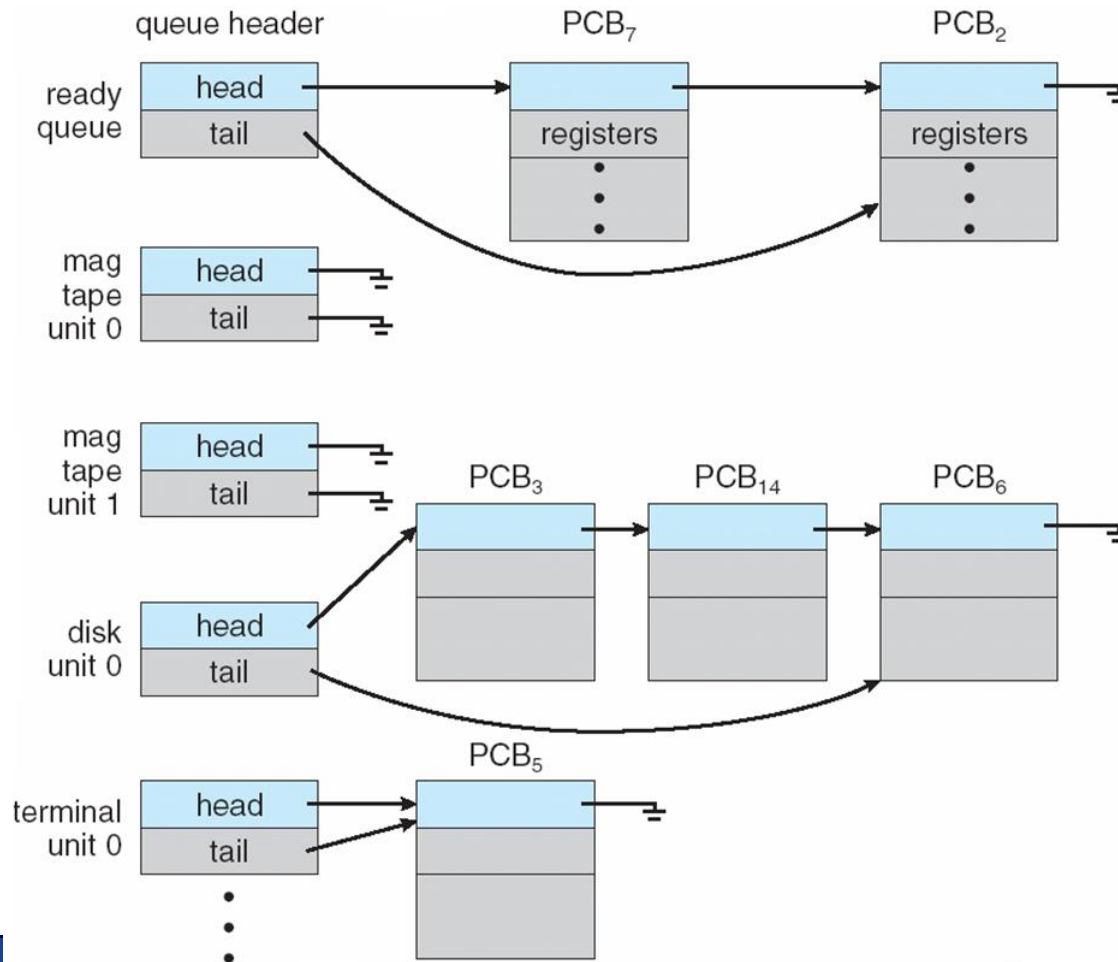
- **Job queue**
 - List of all processes in system

- **Ready queue**
 - List of processes, residing in main memory, ready to execute

- **Device queue**
 - List of processes waiting for a particular I/O device
 - Each device has its own device queue

Scheduling Queue


- Each queue is usually represented by linked list



Queueing Diagram


■ Representation of process scheduling

- A process migrates among various scheduling queues throughout its lifetime




Schedulers

- Scheduler selects processes from queues in some fashion
 - Long-term scheduler (job scheduler)
 - Short-term scheduler (CPU scheduler)



Schedulers

- Short-term scheduler (CPU scheduler)



- Executed frequently (at least once every 100 msec.)
- Scheduling time should be very short

Schedulers

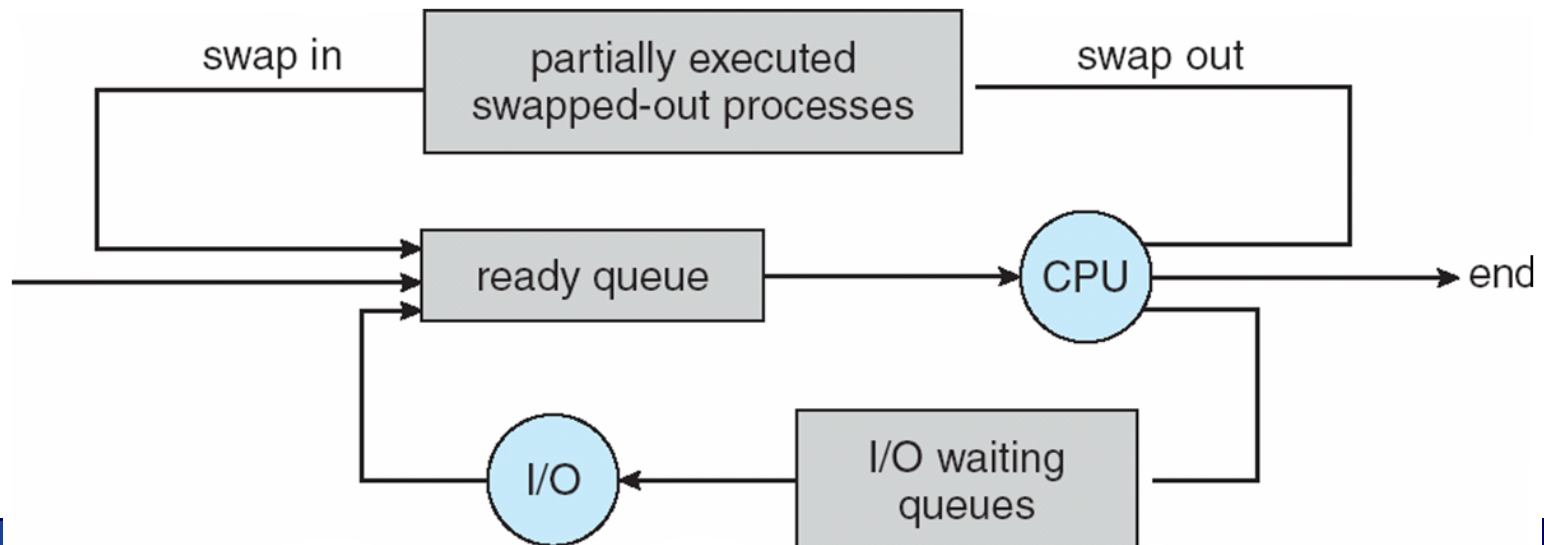
■ Long-term scheduler (job scheduler)



- Controls **degree of multiprogramming**
 - In stable state, average process creation rate == average process departure rate
- Executed less frequently
 - Executed only when a process leaves the system
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Hopefully, long-term scheduler should select a good mix of **I/O-bound** and **CPU-bound** processes

Schedulers

- In some systems, long-term scheduler may be absent or minimal
Ex) UNIX, Windows
 - System stability depends on physical limitation or self-adjusting nature of human
- Some time-sharing system has **medium-term scheduler**
 - Reduce degree of multiprogramming by removing processes from memory



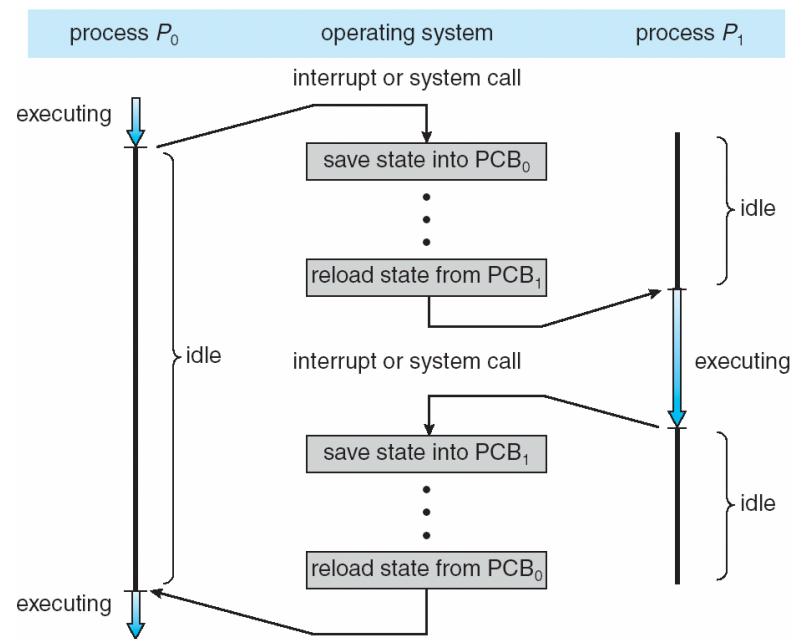

Multitasking in Mobile Systems



- **Some systems / early systems allow only one process to run, others suspended**
- **Due to screen real estate, user interface limits iOS provides for a**
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- **Android runs foreground and background, with fewer limits**
 - Background process uses a **service** to perform tasks
 - **A service**: a separate application component that runs on behalf of the background process.
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use

Context Switch

- Switching running process requires **context switch**
 - State save of current process (PCB)
 - State restore of different process



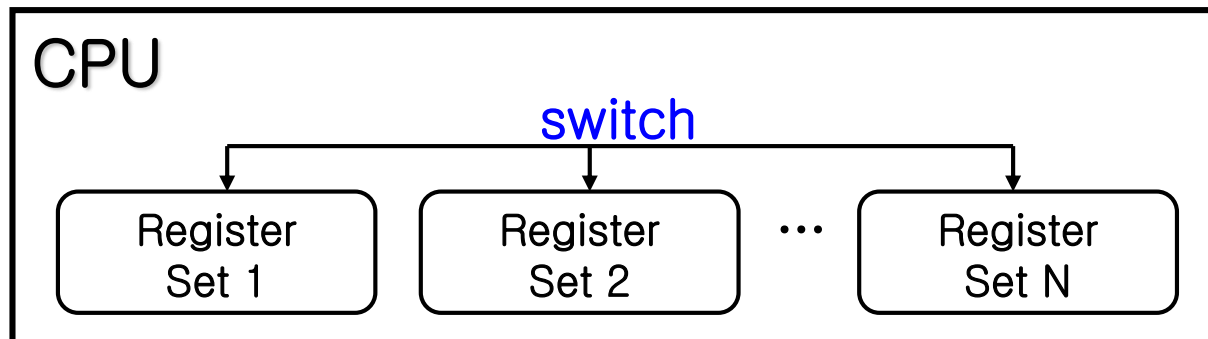
Context Switch



- **Context switch**: the computing process of storing and restoring the state (context) of a CPU such that multiple processes can share a single CPU resource.
- “Context” includes
 - Register contents
 - OS specific data
 - Extra data required by advanced memory-management technique
Ex) page table, segment table, ...
- When to switch?
 - Multitasking
 - Interrupt handling

Context Switch

- Context switching requires considerable overhead.
- H/W supports for context-switching
 - H/W switching (single instruction to load/save all registers)
cf. However, S/W switching can be more selective and save only that portion that actually needs to be saved and reloaded.
 - Multiple set of register for fast switching
Ex) UltraSPARC



Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

Operations on Processes




- Process create
- Process termination
- Process communication

Process Creation

- Create-process system call

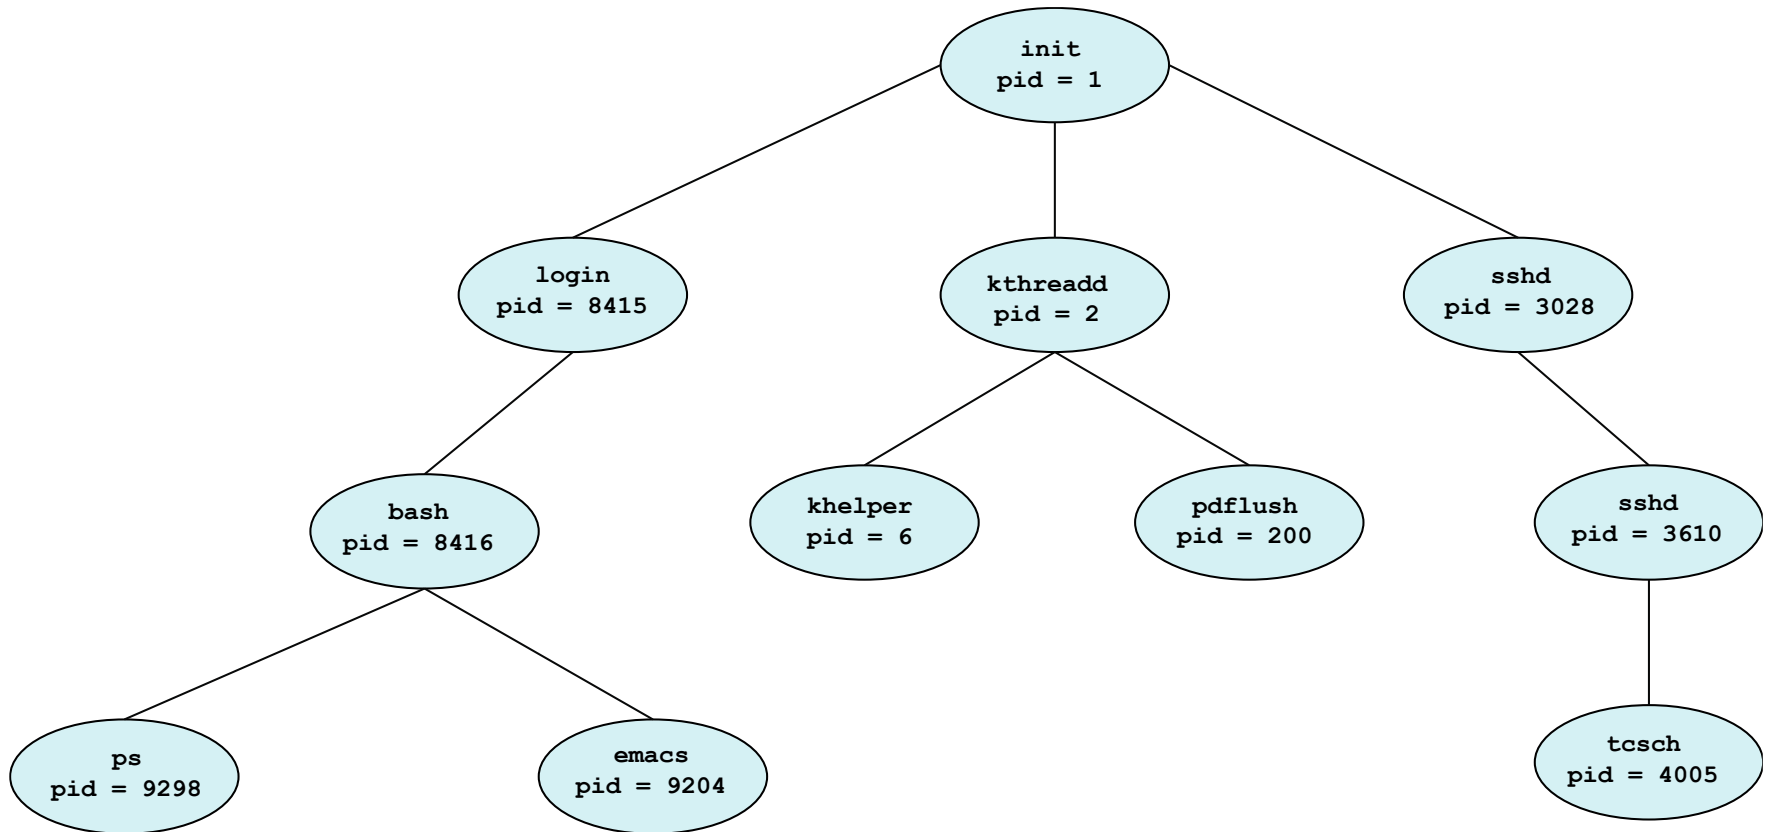
- Create a process and assign a pid.



- Process tree

- Parent-child relation between processes

A Tree of Processes in Linux



Displaying Process Information



- UNIX

- ps [-el]

- Windows

- Task manager (windows system program)
 - Process explorer (freeware)
 - Downloadable from FTP server

Process Creation

■ Some options to create a process

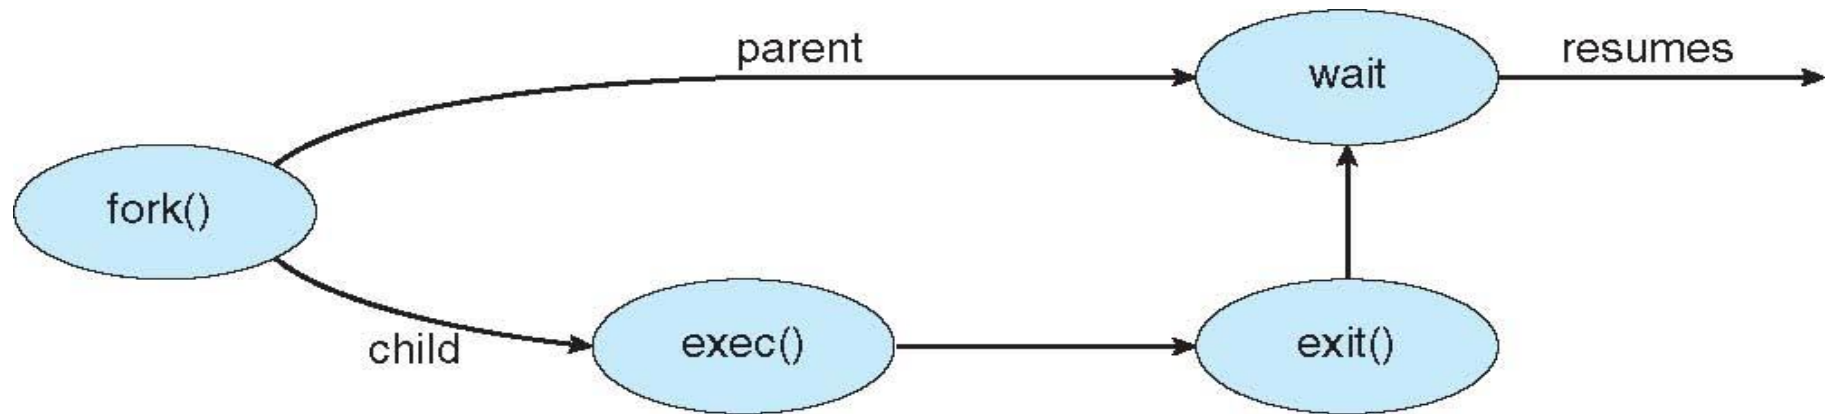
	Options
Resource	<ol style="list-style-type: none">1. Parent and children share all resources2. Children share subset of parent's resources3. Parent and child share no resources
Execution	<ol style="list-style-type: none">1. Concurrent execution2. Parent waits until child is terminated
Address space	<ol style="list-style-type: none">1. Child duplicate of parent2. Child has a program loaded into it

Process Creation in UNIX



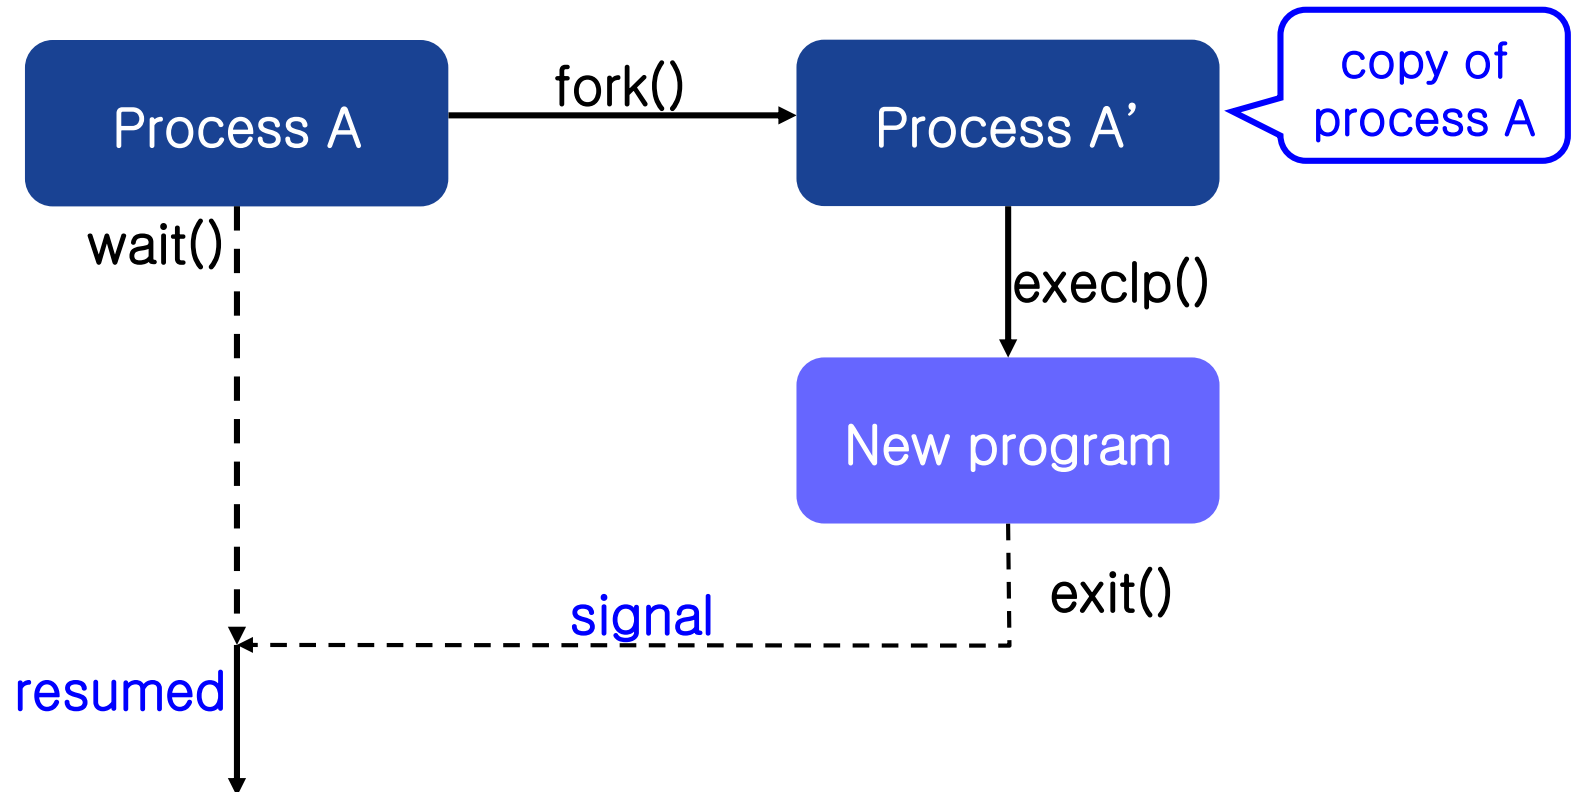
- UNIX system calls related to process creation
 - **fork()**: create process and returns its pid
 - In parent process, return value is **pid of child**
 - In child process, return value is zero
 - **exec() family**: execute a program. The new program substitutes the original one.
 - `execl()`, `execv()`, `execvp()`, `execle()`, `execve()`
 - **wait()**: waits until child process is terminated

Process Creation



Example of Process Creation

- Executing other program



Example of Process Creation

```
int main()
{
    pid_t pid = fork();    // create a process
    if(pid < 0){            // error occurred
        fprintf(stderr, "fork failed\n");
        exit(-1);
    } else if(pid == 0){    // child process
        execlp("/bin/ls", "ls", NULL);
    } else {                // if pid != 0, parent process
        wait(NULL);         // waits for child process to complete
        printf("Child Completed\n");
        exit(0);
    }
}
```


Example of Process Creation

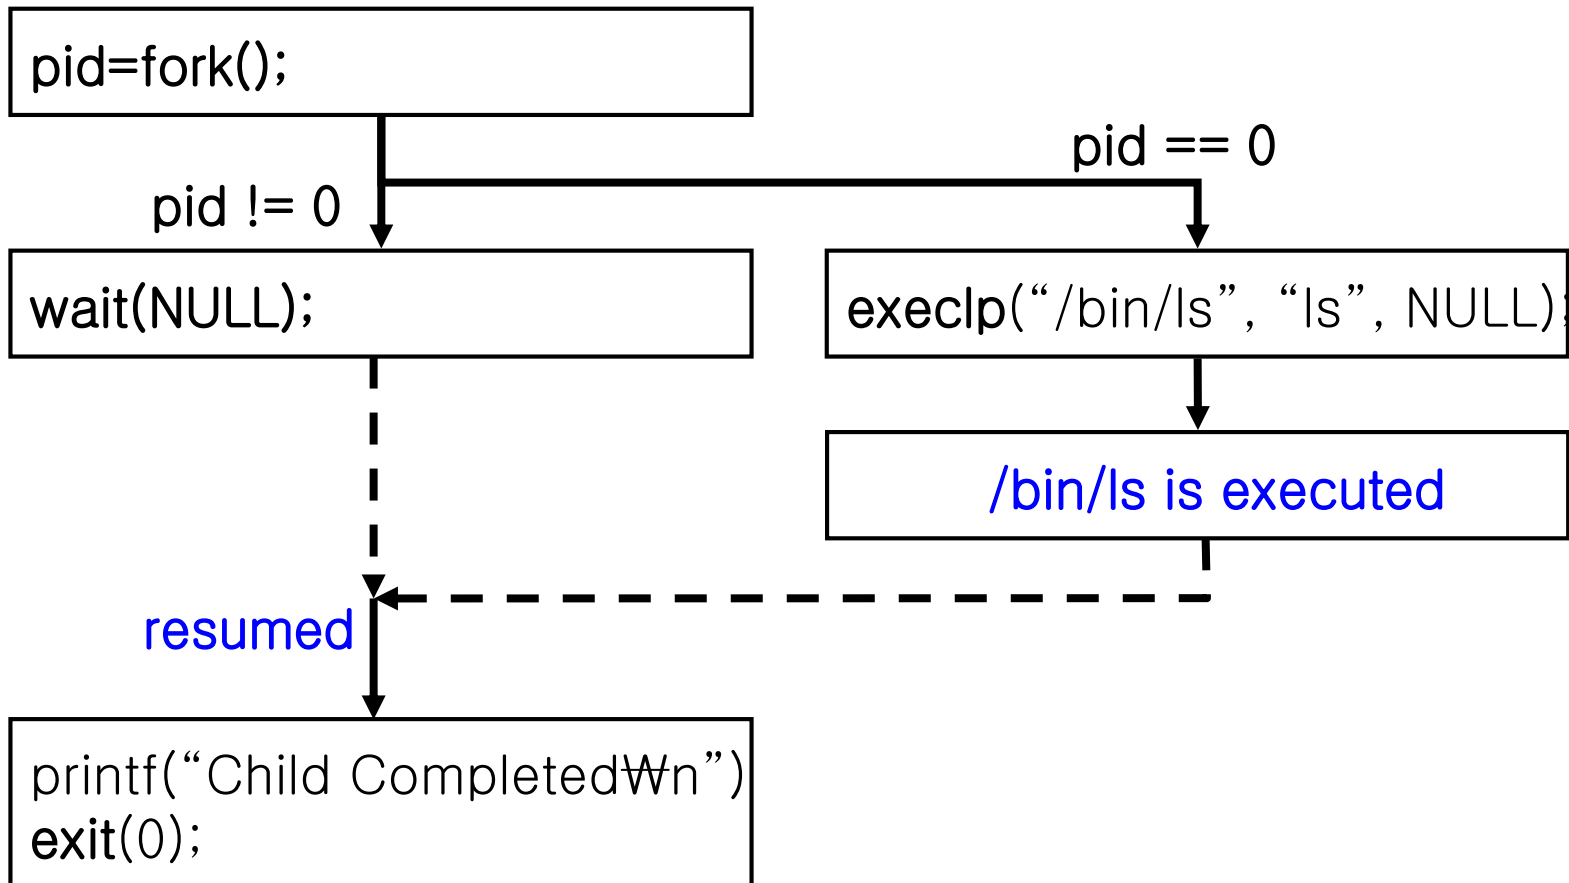
■ Parent process

```
int main()
{
    pid_t pid = fork();
    if(pid < 0){
        fprintf(stderr, "fork failed\n");
        exit(-1);
    } else if(pid == 0){
        execlp("/bin/ls", "ls", NULL);
    } else {
        wait(NULL);
        printf("Child Completed\n")
        exit(0);
    }
}
```

■ Child process

```
int main()
{
    pid_t pid = fork();
    if(pid < 0){
        fprintf(stderr, "fork failed\n");
        exit(-1);
    } else if(pid == 0){
        execlp("/bin/ls", "ls", NULL);
    } else {
        wait(NULL);
        printf("Child Completed\n")
        exit(0);
    }
}
```


Example of Process Creation



More About fork()

■ Resource of child process

- Data (variables): copies of variables of parent process
 - Child process has its own address space
 - The only difference is **pid** returned from **fork()**
- Files
 - Opened before **fork()**: shared with parent
 - Opened after **fork()**: not shared



Process Creation in win32



- **CreateProcess()**
 - Similar to fork() of UNIX, but much more parameters to specify properties of child process
- **WaitForSingleObject()**
 - Similar to wait() of UNIX
- **void ZeroMemory(PVOID *Destination*, SIZE_T *Length*);**
 - Fills a block of memory with zeros.

For more detail, please refer MSDN homepage
(<http://msdn.microsoft.com>)

System Calls in win32

■ Process creation

```
BOOL CreateProcess(  
    LPCTSTR lpApplicationName,  
    LPTSTR lpCommandLine,  
    LPSECURITY_ATTRIBUTES lpProcessAttributes,  
    LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    BOOL bInheritHandles,  
    DWORD dwCreationFlags,  
    LPVOID lpEnvironment,  
    LPCTSTR lpCurrentDirectory,  
    LPSTARTUPINFO lpStartupInfo,  
    LPPROCESS_INFORMATION lpProcessInformation  
);
```

System Calls in win32



■ Wait

```
DWORD WaitForSingleObject(  
    HANDLE hHandle,          // pi.hProcess  
    DWORD dwMilliseconds  
);
```

Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\\\WINDOWS\\\\system32\\\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Process Termination

■ Normal termination

- `exit(int return_code)`: invoked by child process

- Clean-up actions
 - Deallocate memory
 - Close files
 - ETC.

- `return_code` is passed to parent process

- Usually, 0 means success
- Parent can read the return code

```
int status = 0;
wait(&status);           // wait until the child is terminated.
ret = WEXITSTATUS(status); // return_code from the child
```


Abnormal Termination



- Possible reasons of abnormal termination
 - Child has exceeded allocated resource.
 - The task assigned child process is no longer required.
 - Parent is terminated, and OS doesn't allow child process to continue (ex: VMS).
 - Cascading termination
- System calls for process termination
 - Self-termination
 - Ex) abort() of UNIX: Send SIGABRT signal to OS to make core dump.
 - Terminating other process
 - Ex) TerminateProcess() of win32

Usually, such a system call can be invoked only by its parent process.