

3. Process Concept

[ECE321/ITP302] Operating Systems

Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

Inter-process Communication (IPC)



■ Goal of IPC: cooperation

- Information sharing
 - Shared file, ...
- Computation speedup
 - Multiple CPU or I/O
- Modularity
 - Dividing system functions
- Convenience
 - Editing, printing, compiling in parallel

■ IPC Models

- Shared-memory model
- Message-passing model

Shared-Memory Systems

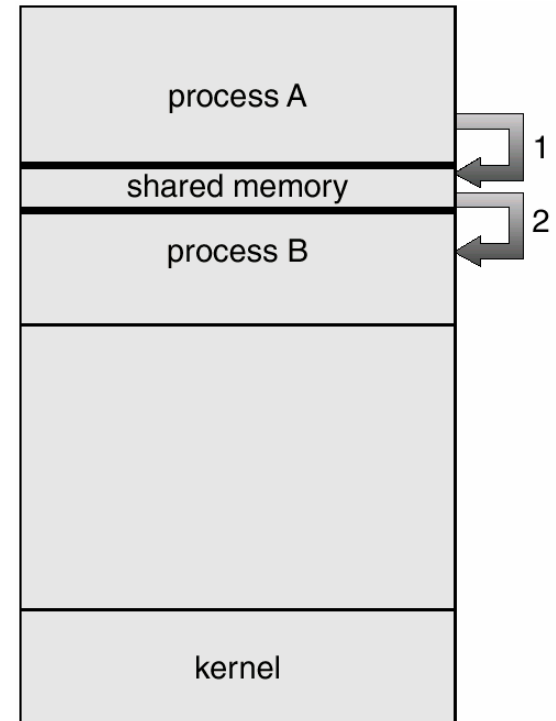
■ Shared-memory segment

- Special memory space that can be shared by two or more processes
- Form of data and location is not determined by OS, but those processes
 - Processes should avoid simultaneous writing by themselves

■ Advantage


- Fast
 - Suitable for large amount of data

Example) producer-consumer problem



Producer–Consumer Problem

- Producer and consumer communicate information (item) through shared memory


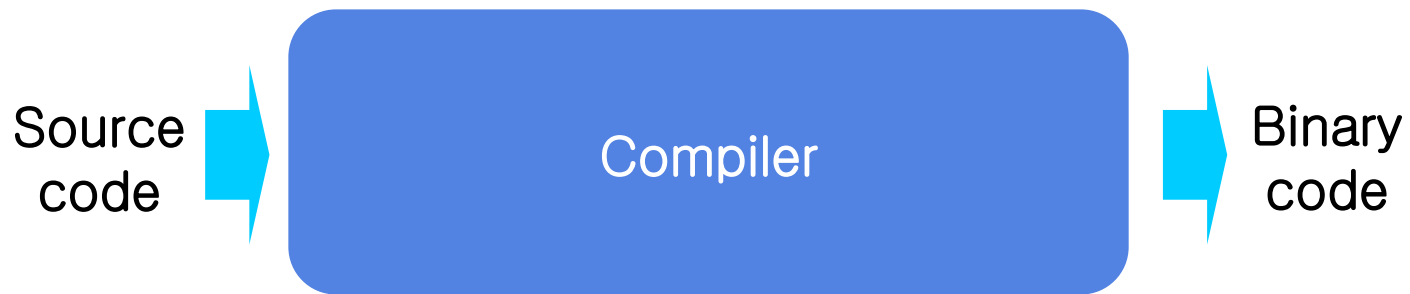


- **Producer**: produce information for consumer
 - **Consumer**: consume information written by producer
- Ex) compiler – assembler, server – client

Note! Producer and consumer should be synchronized.

→ Discussed in chapter 6


Producer-Consumer Problem



Producer–Consumer Problem

■ Two types of buffer

- Unbounded buffer
 - No practical limit on buffer size
 - Producer can always produce
- Bounded buffer
 - Producer must wait if buffer is full.




Producer–Consumer Problem using Bounded Buffer

■ Representation of buffer

- Buffer is represented by **circular queue**

```
#define BUFFER_SIZE 6
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;           // tail or rear
int out = 0;          // head or front
```



■ Empty/full condition

- $in == out$: buffer is empty
 - $(in+1) \% BUFFER_SIZE == out$: buffer is full
- Cf. Buffer can store at most $BUFFER_SIZE - 1$ items

Producer–Consumer Problem using Bounded Buffer

■ Producer


item nextProduced;

```
while (1) {  
    // produce an item in nextProduced  
    while (((in + 1) % BUFFER_SIZE) == out); // do nothing  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

■ Consumer

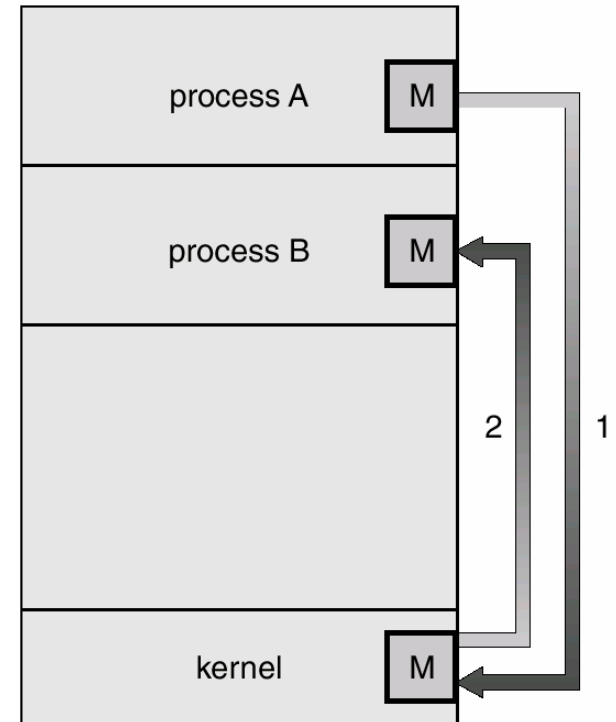
item nextConsumed;

```
while (1) {  
    while (in == out); /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    // consume the item in nextConsumed  
}
```



Message-Passing Systems

- Process communication via passage-passing facility provided by OS
- Advantage
 - No conflict
 - Suitable for smaller amounts of data
 - Communication between processes on different computer




Message-Passing Systems




- For message passing, communication link should be exist between the processes
- Essential operations
 - send(message)
 - receive(message)
- (Logical) Implementation methods
 - Direct/indirect
 - Synchronous/asynchronous
 - Buffering
 - Zero/bounded/unbounded capacity

Direct/Indirect Communication

- Direct communication: connection link directly connects processes



- Indirect communication: processes are connected via mailbox



Direct Communication




- Processes are connected directly
- Symmetry vs. asymmetry in addressing
 - Symmetric addressing: both sender and receiver know each other
 - Sender: `send(P, message)`
 - Receiver: `receive(Q, message)`
 - Asymmetric addressing: only sender knows receiver
 - Sender: `send(P, message)`
 - Receiver: `receive(id, message)`
 - `id` is set to name of sender (output argument)

Direct Communication

■ Properties

- The processes know each other.
- A link is associated with exactly two processes
- Between a pair of processes, only one link exists



■ Disadvantage: limited modularity

- Changing identifier of a process requires examining all other process definitions


Indirect Communication

- Processes are connected indirectly through a mailbox or port

- Sender: `send(A, message)`
- Receive: `receive(A, message)`

- Who is the owner of mailbox?

- Owned by a process
 - No confusion on receiver
- Owned by OS
 - Ownership can be transferred



- Properties

- Processes are connected if they share a common mailbox
- A link may be associated with more than two processes
- Between a pair of processes, there may be a number of different links

Indirect Communication

■ How to avoid confliction?



■ Solutions

- Allow a link to be associated with two processes at most
- Allow at most one process at a time to execute receive()
- Allow the system to select a process to receive arbitrary


Synchronization

■ Sender

- Blocking send: sender is blocked until receiver takes message
- Non-blocking send: sender just send message and resumes operation

■ Receiver

- Blocking receive: if message is not available, receiver waits
- Non-blocking receive: if message is not available, receiver is not blocked but receives a null message



Synchronization



- Different combinations possible
 - If both send and receive are blocking, we have a rendezvous
- Producer–consumer becomes trivial

```
message next produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next produced);  
}
```

```
message next consumed;  
while (true) {  
    receive(next consumed);  
  
    /* consume the item in next consumed */  
}
```

Buffering

- During communication, messages are stored in temporary queue (buffer)



- Three kinds of buffer capacity
 - Zero capacity: only blocking send is possible
 - Bounded capacity: buffer has finite length n
 - If buffer is full, sender must be blocked
 - Otherwise, sender can resume
 - Unbounded capacity: buffer has infinite capacity
 - Sender never blocks

Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

Examples of IPC Systems



- Shared-memory (POSIX)
- Message-passing (MACH)
- Local Procedure Call (Windows XP)
 - Undocumented internal API (Skip)

Examples of IPC Systems – POSIX



POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment to share it

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Now the process could write to the shared memory

```
sprintf(shared memory, "Writing to shared memory");
```

IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```

IPC POSIX

Producer



```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```


[Mach] Message Passing



- Mach: an OS kernel developed at CMU
 - One of the earliest examples of a **microkernel**
 - Even system calls are made **by message**
 - Especially designed for multiprocessing, distributed system
 - Execution unit: task (similar to process, but has multiple thread)
- If a task is created, two mailboxes (ports) are also created
 - Kernel mailbox → communication with task
 - Notify mailbox → notification of event occurrences
- Message passing system calls
 - `msg_send()`, `msg_receive()`
 - `msg_rpc()` – Remote Procedure Call that can work between systems.

[Mach] Message Passing



- **Creating mailbox**

- port_allocate(): create new mailbox

- **Mailbox of Mach**

- New mailbox is owned by the process that creates it.
 - At a time, only one process can own and receive from a mailbox
 - Basically, FIFO order
 - But order of messages from other processes are not guaranteed

- **Mailbox set**

- A set of mailboxes which is treated as single mailbox
 - Each mailbox is assigned to a thread in a task

- **Major problem: poor performance**

- Double copying of messages (sender → mailbox → receiver)
 - For local communication, a remedy is provided

[Mach] Message Passing



- Send and receive are flexible, for example four options if mailbox full:
 - Wait indefinitely
 - Wait at most n milliseconds
 - Return immediately
 - Temporarily cache a message

[Windows] Message Passing

- Message-passing centric via (Advanced) local procedure call (LPC) facility
 - Only works between processes on the same system (machine)
 - Similar to RPC, but optimized for WindowsXP
 - Uses ports (like mailboxes) to establish and maintain communication channels
- Communication works as follows:
 - The client opens a handle to the subsystem's connection port object
 - The client sends a connection request
 - The server creates two private communication ports and returns the handle to one of them to the client
 - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies

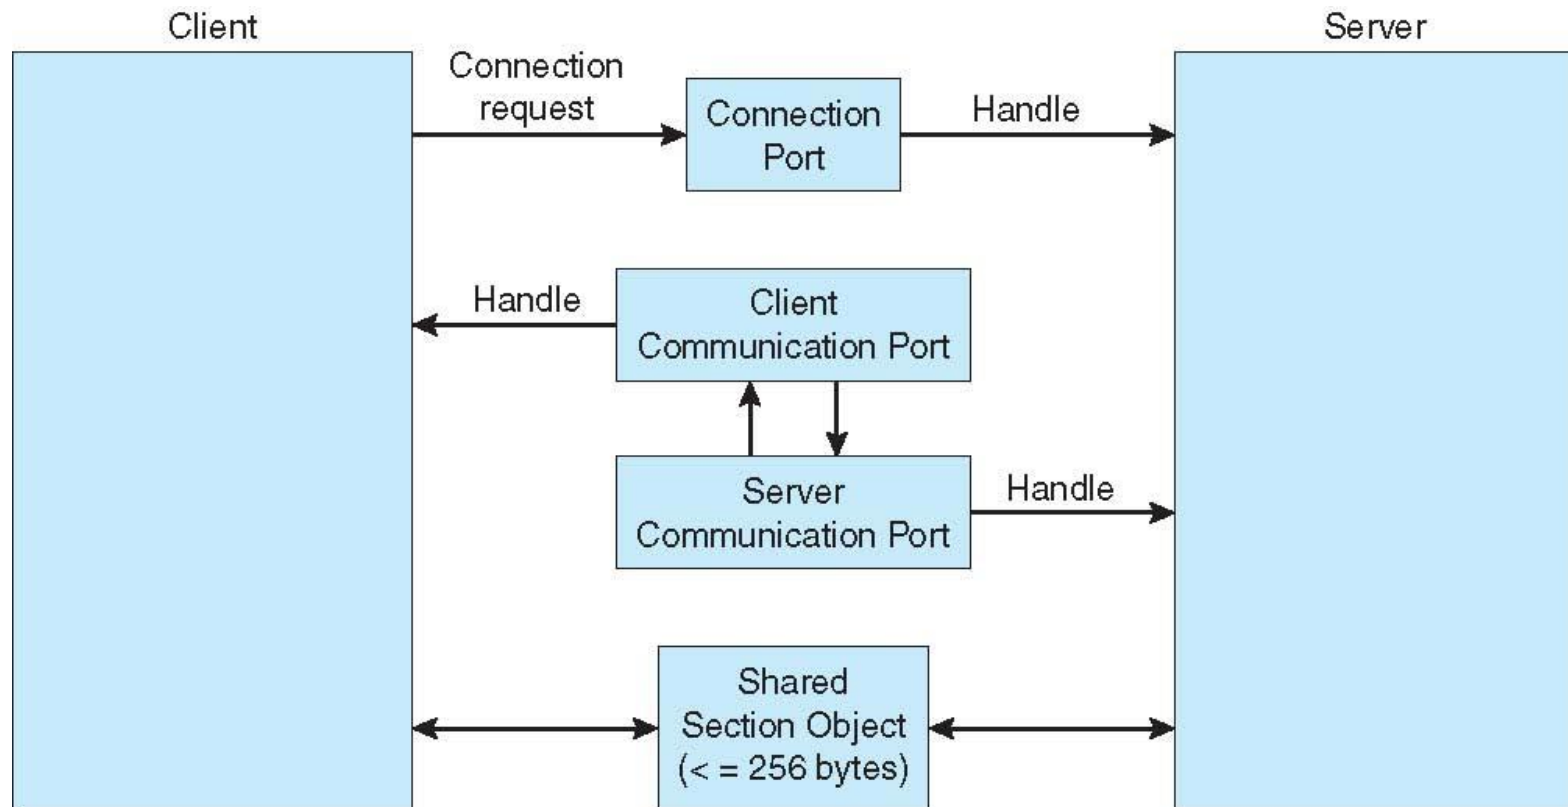
[Windows] Message Passing



- Three possible message-passing techniques
 - 1) If message is small (≤ 256 bytes), the port's message queue is used as the intermediate storage
 - 2) Other wise, the message is passed through section object which is a region of shared memory.
(Avoids data copying)
 - 3) Callback mechanism can be used, when client or server cannot respond immediately.

- LPC is not a part of win32
 - LPC facility is not visible to the application programmer
 - Applications invoke standard RPC.
 - If RPC is being invoked on a process on the same system, RPC is indirectly handled through LPC

ALPC in windows




Agenda



- Overview
- Process scheduling
- Operations on processes
- Inter-process communication
- Example of IPC system
- Communication in client-server systems

Client-Server



Communications in Client–Server Systems



- **Socket**
 - Data communication

- **RPC (Remote Procedure Call)**
 - Procedure call between systems (machines)
 - Procedural programming


- **RMI (Remote Method Invocation) of JAVA**
 - Invoking method of object in other system
 - Object oriented programming

Socket

- **Socket**: logical endpoint for communication




- Identified by <ip address>:<port #>



- Each connection is identified by a pair of sockets


Socket

- **Port**: logical contact point to a computer recognized by TCP and UDP protocols
 - A computer may have multiple ports (0 ~ 65535)



- Well-known services have their own ports below 1024
Ex) telnet: 23, ftp: 21, http: 80
 - Server always listens corresponding port
- Ports above 1024 can be arbitrary assigned for network communication

Socket



Socket




- Server opens a port to accept connection request
- Initiating connection
 - Client arbitrary assigns a port above 1024
Ex) a client 146.86.5.20 assigned a port 1625
 - Client request a connection to server
Ex) a web server 161.25.19.8 (port # of web service: 80)
 - If server accepts request, connect is established
Ex) <146.86.5.20:1625> – <161.25.19.8:80>

Java Socket

■ Socket classes

- ServerSocket: accepts request for connection
- Socket: in charge of actual communication



```
import java.net.*;
import java.io.*;
public class DateServer{
public static void main(String[] args) {
    try {
        ServerSocket sock = new ServerSocket(6013);

        // now listen for connections
        while (true) {
            Socket client = sock.accept();
            // we have a connection

            PrintWriter pout = new PrintWriter(client.getOutputStream(), true);
            // write the Date to the socket
            pout.println(new java.util.Date().toString());

            // close the socket and resume listening for more connections
            client.close();
        }
    } catch (IOException ioe) {
        System.err.println(ioe);
    }
}}
```

```
import java.net.*;
import java.io.*;
public class DateClient{
public static void main(String[] args) {
    try {
        // this could be changed to an IP name or address other than the localhost
        Socket sock = new Socket("127.0.0.1",6013);
        InputStream in = sock.getInputStream();
        BufferedReader bin = new BufferedReader(new InputStreamReader(in));

        String line;
        while( (line = bin.readLine()) != null)
            System.out.println(line);
        sock.close();
    }
    catch (IOException ioe) {
        System.err.println(ioe);
    }
}}
```

Java Socket

■ Server

1. Create a ServerSocket

```
ServerSocket socket = new  
    ServerSocket(6013);
```

2. Wait for a client

```
Socket client =  
    socket.accept();
```

- 4a. If a client is accepted,
 communicate with client via
 client

■ Client

3. Create a socket to server

```
Socket sock = new  
    Socket("127.0.0.1", 6013);
```

- 4b. If connection was
 established, communicate
 with server via *sock*

Java Socket

■ Server (given *client*)

```
PrintWriter pout = new  
    PrintWriter(client.getOutputS  
        tream(), true);
```

```
pout.println(new  
    java.util.Date().toString());
```

```
client.close();
```

■ Client (given *sock*)


```
InputStream in =  
    sock.getInputStream();  
BufferedReader bin = new  
    Buffered Reader(new  
        InputStreamReader(in))
```

```
String line;  
while((line = bin.readLine()) !=  
    null)  
    System.out.println(line);
```

```
sock.close();
```

Remote Procedure Calls (RPC)

- **RPC**: procedure call mechanism between systems
- On server, **RPC daemon** listens to a port
- Client sends a message containing identifier of function and parameters



Remote Procedure Calls



- RPC is served through stubs
 - Client invoke remote procedure as it would invoke a local procedure call
- **Stub**: a small program providing interface to a larger program or service on remote side
 - Client stub / server stub
 - Locate port on server
 - Marshal / unmarshal parameters

Remote Procedure Calls

■ Parameter marshaling

Motivation: each system has its own data format


Ex) representation of integer on a system may differ from that on other system

→ parameter should be transferred in standard format

□ XDR: eXternal Data Representation

■ **Marshalling**: native representation → XDR

■ **Unmarshalling**: XDR → native representation



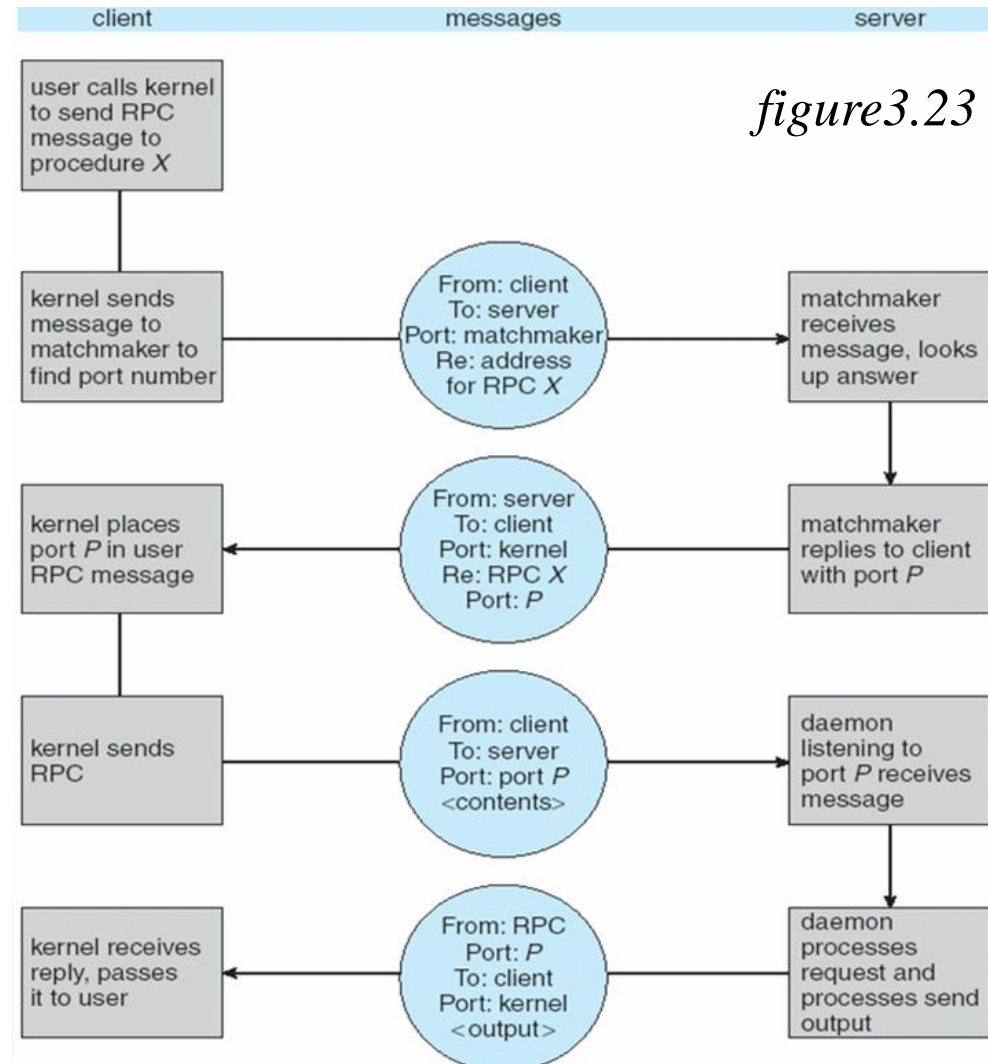
Remote Procedure Calls



- Every RPC request should serviced “exactly once”
 - Attach a timestamp to each message
 - Server keeps history of message it has served
 - At every request, it checks the history.
 - Server send ACK message to client
 - Client resend RPC call periodically, until it receives ACK.

Remote Procedure Calls

- Issue: how to bind the client and the server port with no information at start?
- Two approaches for assigning RPC port
 - 1) Fixed address (hard-coding)
 - 2) Transferred through rendezvous daemon (**matchmaker**) illustrated in the figure3.23



RPC Reference Sites



- Windows

- MSDN RPC page:

- <http://msdn.microsoft.com/library/default.asp?url=/library/en-us/dnanchor/html/rpcank.asp>

- Unix

- Document about *rpcgen*.

Others



■ Pipes

- Acts as a conduit allowing two processes to communicate
- Issues
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e. *parent-child*) between the communicating processes?
 - Can the pipes be used over a network?

■ Ordinary pipes: Unidirectional, parent-child

■ Named Pipes: Bidirectional, no parent-child